

Information Access with Apache Lucene - Part 1

Metodi per il Ritrovamento dell'Informazione

Laurea Triennale in Informatica
Università degli Studi di Bari Aldo Moro

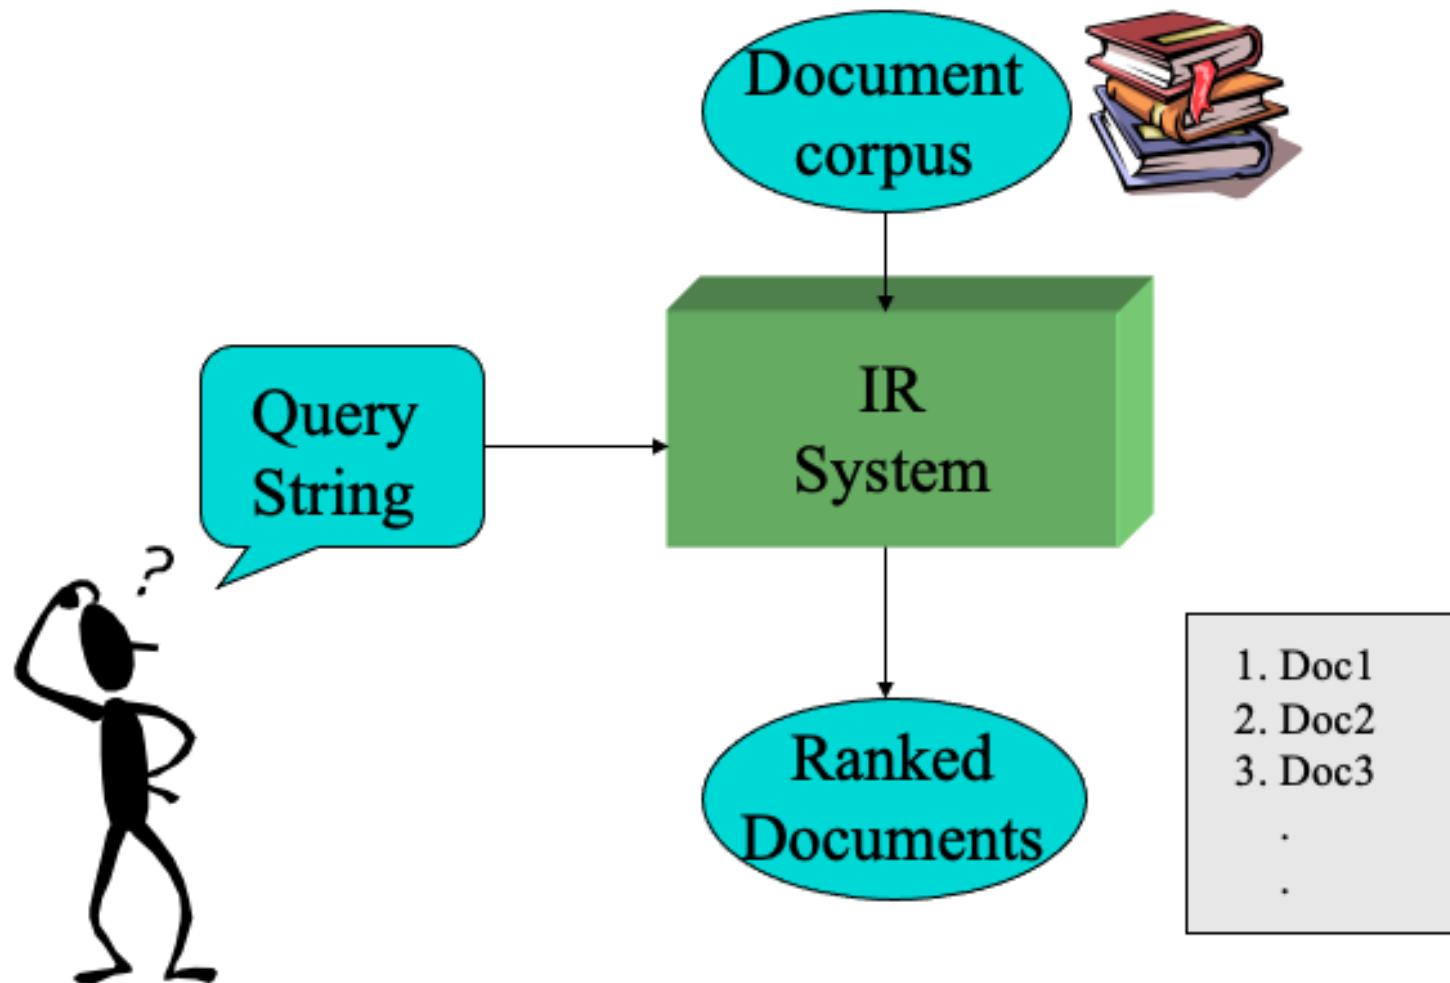
Prof. Cataldo Musto

cataldo.musto@uniba.it

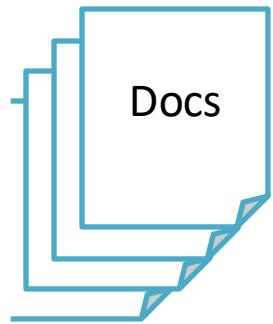
Recap

SEARCH ENGINE

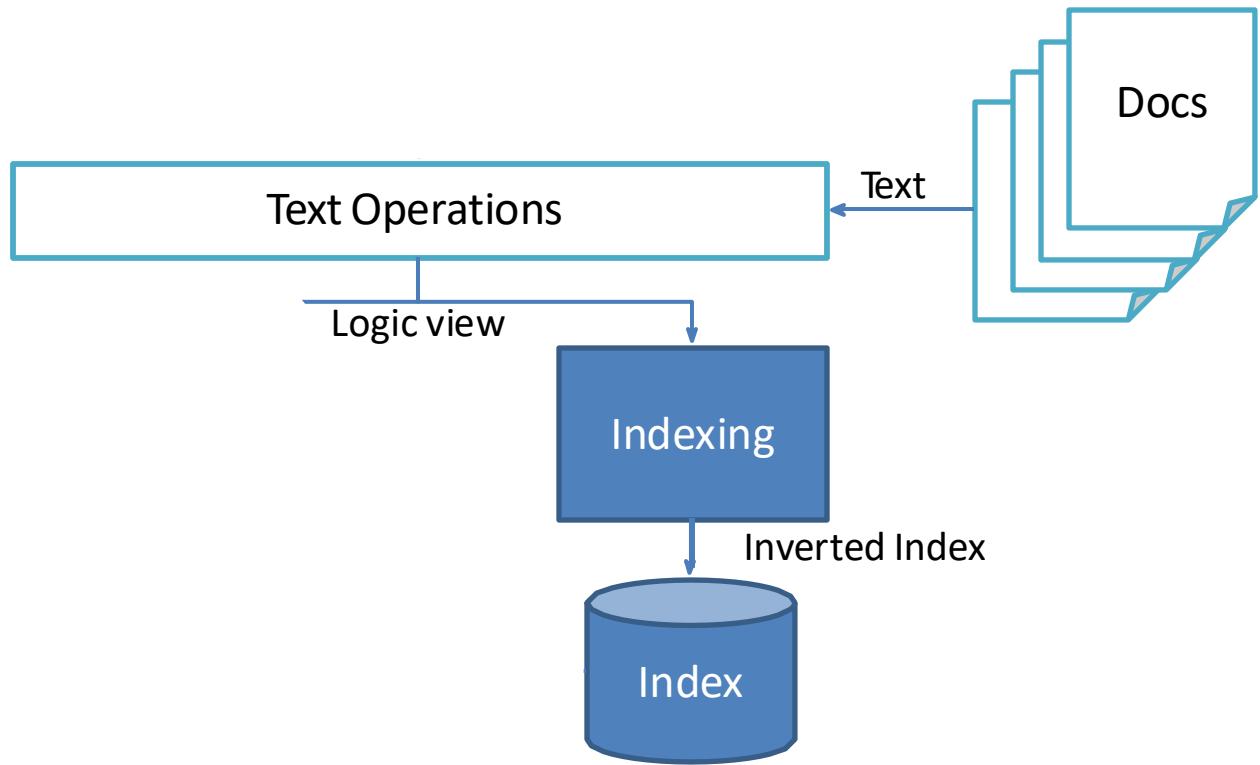
Information Retrieval Process (recap)



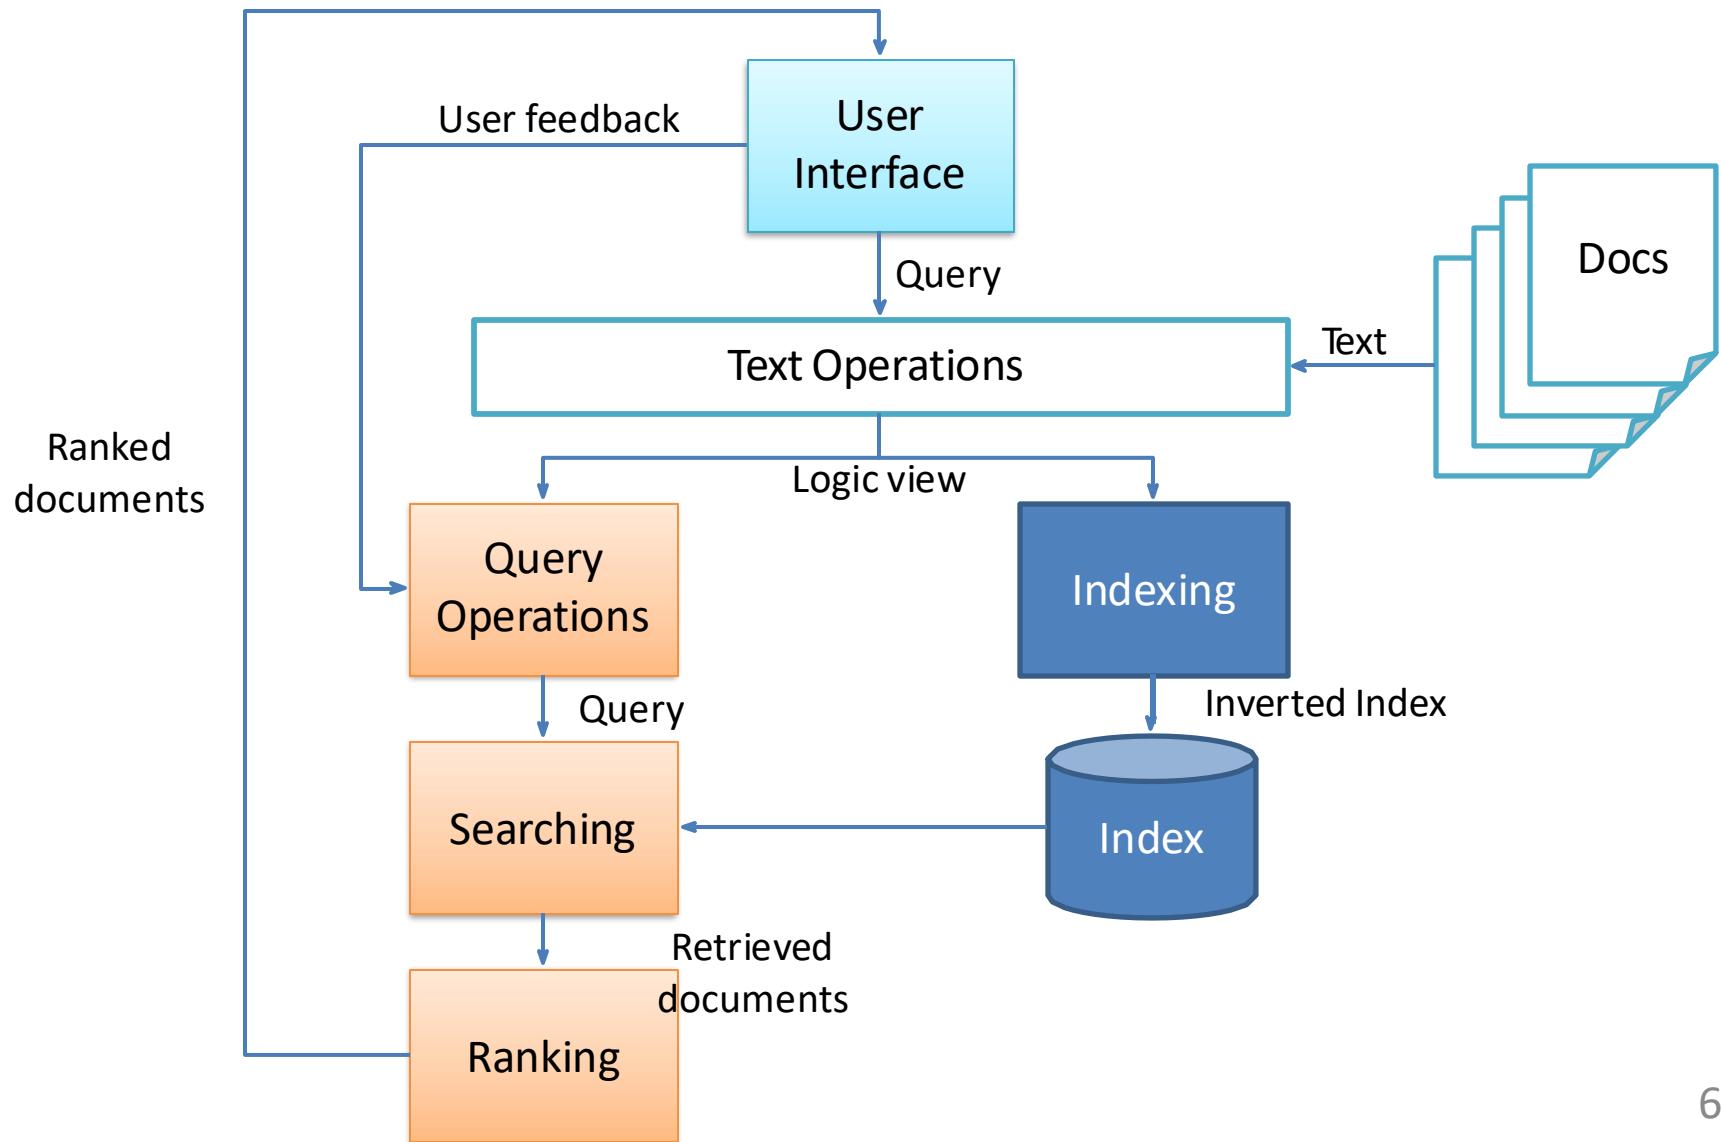
Information Retrieval Process (recap)



Information Retrieval Process (recap)



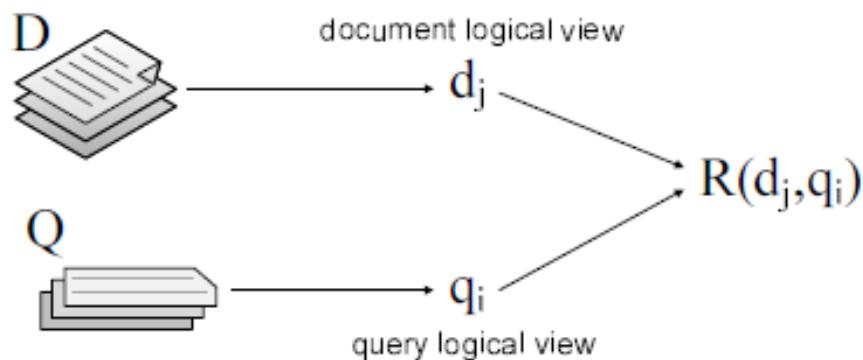
Information Retrieval Process (recap)



Information Retrieval Model

$\langle D, Q, F, R(q_i, d_j) \rangle$

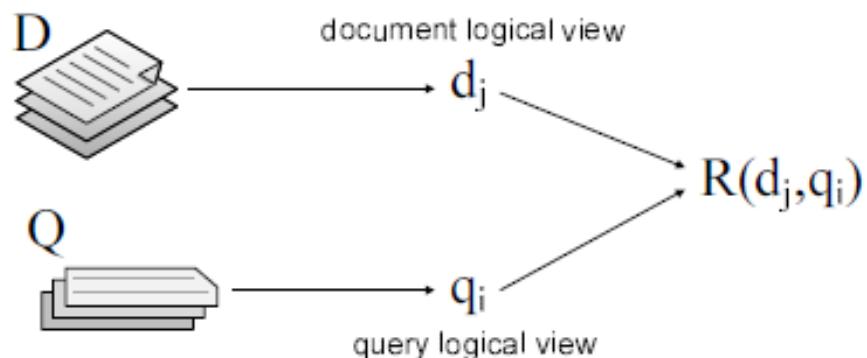
- D: document
- Q: query
- F: query/document representation function
- $R(q_i, d_j)$: ranking function



Information Retrieval Model

$\langle D, Q, F, R(q_i, d_j) \rangle$

- D: document
- Q: query
- F: query/document representation function ?
- $R(q_i, d_j)$: ranking function



Bag-of-words representation

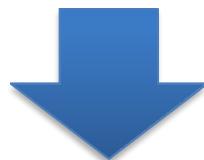
Document/query as unordered collection of words

John likes to watch movies. Mary likes too. John also likes to watch football games.

Bag-of-words representation

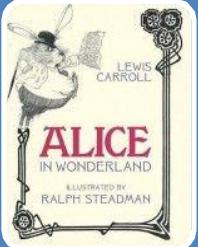
Document/query as unordered collection of words

John likes to watch movies. Mary likes too. John also likes to watch football games.

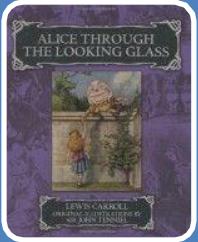


```
{"John": 1, "likes": 2, "to": 3, "watch": 4, "movies": 5, "also": 6, "football": 7, "games": 8, "Mary": 9, "too": 10}
```

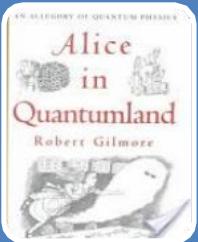
Term-Document matrix



- 'It's a friend of mine — a Cheshire Cat,' said Alice: 'allow me to introduce it.'
- 'It's the oldest rule in the book,' said the King. 'Then it ought to be Number One,' said Alice.

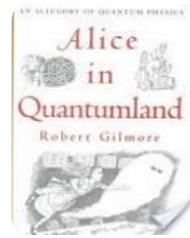
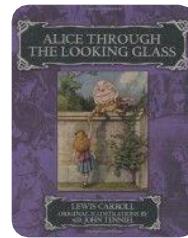
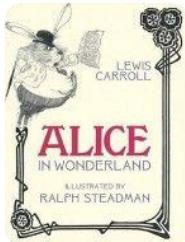


- Alice watched the White King as he slowly struggled up from bar to bar, till at last she said, 'Why, you'll be hours and hours getting to the table, at that rate.'
- Alice looked round eagerly, and found that it was the Red Queen. 'She's grown a good deal!' was her first remark.



- In the pool of light was a billiards table, with two figures moving around it. Alice walked toward them, and as she approached they turned to look at her.
- Alice lay back, and closed her eyes. There was the Red Queen again, with that incessant grin. Or was it the Cheshire cat's grin?

Term-Document matrix



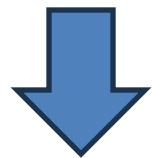
	D1	D2	D3
Cheshire Cat	1	0	1
Alice	2	2	2
book	1	0	0
King	1	1	0
table	0	1	1
Queen	0	1	1
grin	0	0	2

Term-Document matrix

	D1	D2	D3
Cheshire Cat	1	0	1
Alice	2	2	2
book	1	0	0
King	1	1	0
table	0	1	1
Queen	0	1	1
grin	0	0	2

Query: **Alice AND Queen**

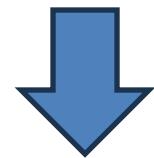
Term-Document matrix



	D1	D2	D3	Q
Cheshire Cat	1	0	1	0
Alice	2	2	2	1
book	1	0	0	0
King	1	1	0	0
table	0	1	1	0
Queen	0	1	1	1
grin	0	0	2	0

Query: **Alice AND Queen**

Term-Document matrix



	D1	D2	D3	Q
Cheshire Cat	1	0	1	0
Alice	2	2	2	1
book	1	0	0	0
King	1	1	0	0
table	0	1	1	0
Queen	0	1	1	1
grin	0	0	2	0

Query: **Alice AND Queen**

Result: D2, D3

This is the representation
that we adopt for
information retrieval tasks

Inverted Index

Index

Page numbers in **bold face** refer to key term definitions

Page numbers in *italics* refer to images or diagrams

Page numbers followed by a "t" indicate a table

A

absolute temperature scale, 350–351

absolute zero, 351

acceleration of gravity, A.23t

accuracy, A.5

acetic acid (CH_3COOH)

 buffers, 575–576, 581–582

 conjugate acid-base pairs, 540

 ionization constant, 553, 554t

 manufacture of, 451

 titrations, 590–592

 as weak acid, 144t, 145, 551–552

acid-base pairs, conjugate, 540–544

acid-base reactions, 538

 autoionization of water, 545–547

 gas-forming exchange, 150–151

 net ionic equations for, 148–150

 neutralization, 146–150, 561–566t

 of salts 146–151 561–566

air, 342–343, 366–370, 380–381, 706

alkyl groups, 70–71

alcohols, 64, 505–507

aldehydes, 278–279

alkali metals, 55, 106

alkaline batteries, 670

alkaline earth metals, 55

alkaline fuel cells, 674

alkalosis, 576

alkanes, 68–71, 277–278, A.25–A.26

alkenes, 280–283, A.26–A.27

alkyl group, 70–71, A.25–A.26

alkynes, 281, A.27

allotropes, 23–24, 208, 403–405

alpha particles, 38–39, 693–696, 697, 699–700

alpha radiation, 693

alpha rays, 36–37

aluminum (Al), 7, 8t, 103, 634–635, 682

amines, 544–545

ammonia (NH_3)

 amines, 545

 Brensted-Lowry base, 538–539

 complex ions, 567

Inverted Index

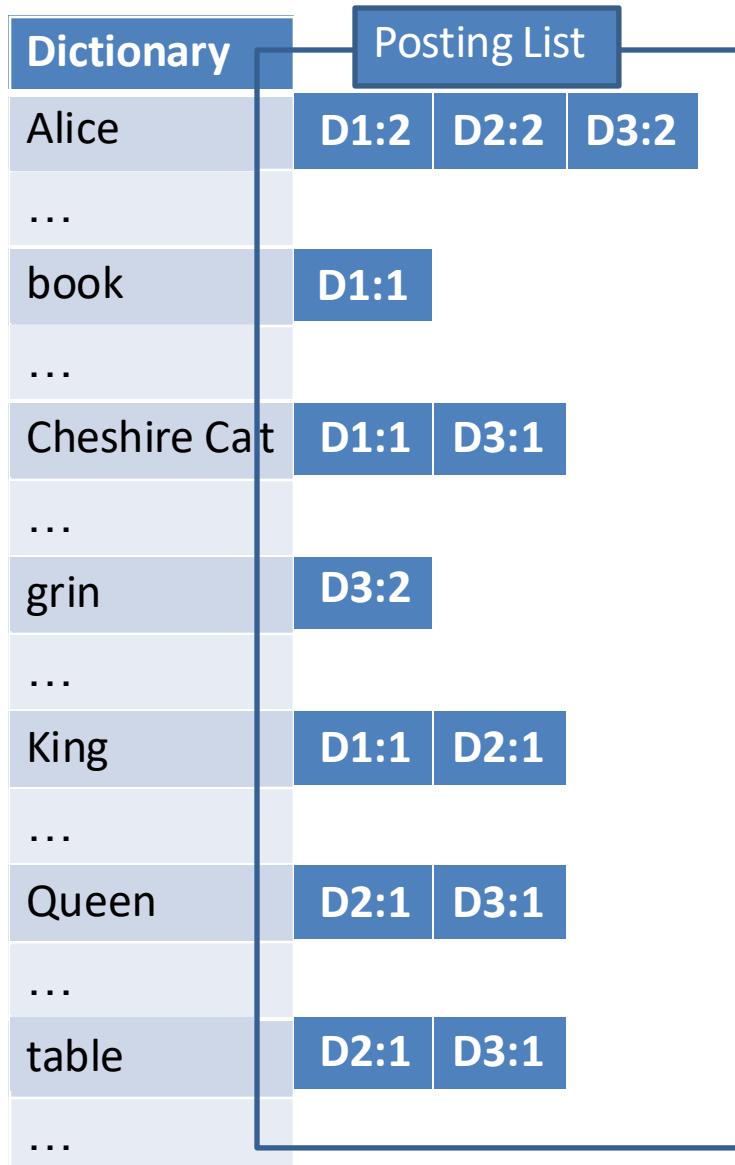
Dictionary	D1:2	D2:2	D3:2
Alice	D1:2	D2:2	D3:2
...			
book	D1:1		
...			
Cheshire Cat	D1:1	D3:1	
...			
grin	D3:2		
...			
King	D1:1	D2:1	
...			
Queen	D2:1	D3:1	
...			
table	D2:1	D3:1	
...			

Inverted Index

Dictionary	D1:2	D2:2	D3:2
Alice	D1:2	D2:2	D3:2
...			
book	D1:1		
...			
Cheshire Cat	D1:1	D3:1	
...			
grin	D3:2		
...			
King	D1:1	D2:1	
...			
Queen	D2:1	D3:1	
...			
table	D2:1	D3:1	
...			

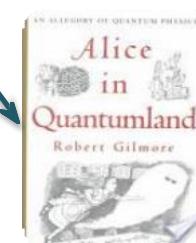
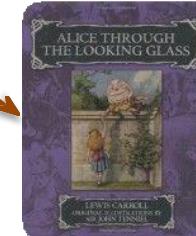
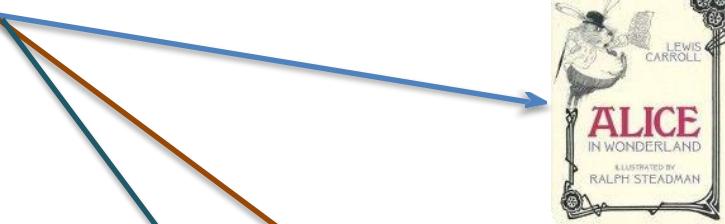
document:occurrences

Inverted Index



Inverted Index

Dictionary	D1:2	D2:2	D3:2
Alice	D1:2	D2:2	D3:2
...			
book	D1:1		
...			
Cheshire Cat	D1:1	D3:1	
...			
grin	D3:2		
...			
King	D1:1	D2:1	
...			
Queen	D2:1	D3:1	
...			
table	D2:1	D3:1	
...			



Inverted Index

Dictionary	D1:2	D2:2	D3:2
Alice	D1:2	D2:2	D3:2
...			
book	D1:1		
...			
Cheshire Cat	D1:1	D3:1	
...			
grin	D3:2		
...			
King	D1:1	D2:1	
...			
Queen	D2:1	D3:1	
...			
table	D2:1	D3:1	
...			

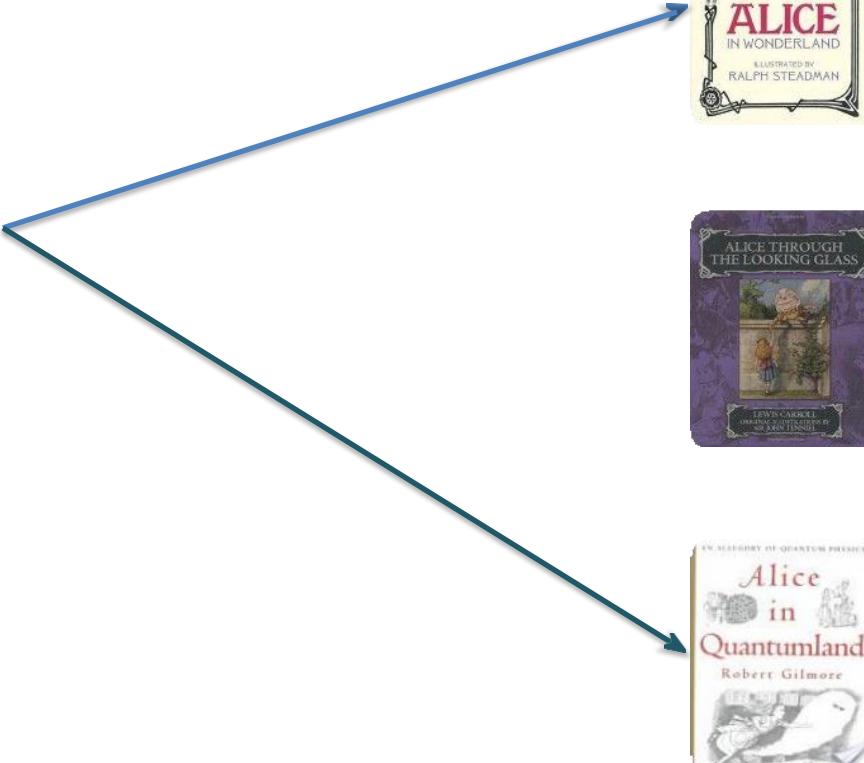
A blue arrow points from the 'book' entry in the dictionary to the book cover of 'Alice in Wonderland' by Lewis Carroll, illustrated by Ralph Steadman.

A blue arrow points from the 'grin' entry in the dictionary to the book cover of 'Alice Through the Looking Glass' by Lewis Carroll, original illustrations by Sir John Tenniel.

A blue arrow points from the 'table' entry in the dictionary to the book cover of 'Alice in Quantumland: An Allegory of Quantum Physics' by Robert Gilmore.

Inverted Index

Dictionary	D1:2	D2:2	D3:2
Alice	D1:2	D2:2	D3:2
...			
book	D1:1		
...			
Cheshire Cat	D1:1	D3:1	
...			
grin	D3:2		
...			
King	D1:1	D2:1	
...			
Queen	D2:1	D3:1	
...			
table	D2:1	D3:1	
...			



The diagram illustrates an inverted index for words found in three books. The words and their document IDs are:

- Alice: D1:2, D2:2, D3:2
- book: D1:1
- Cheshire Cat: D1:1, D3:1
- grin: D3:2
- King: D1:1, D2:1
- Queen: D2:1, D3:1
- table: D2:1, D3:1

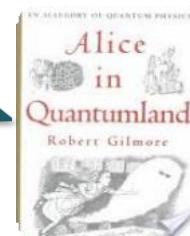
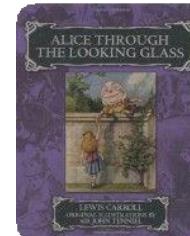
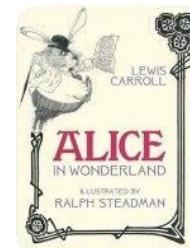
Each word entry is connected by a blue arrow to its respective book cover image:

- Alice: Alice in Wonderland by Lewis Carroll
- book: Alice Through the Looking Glass by Lewis Carroll
- Cheshire Cat: Alice in Quantumland by Robert Gilmore

Inverted Index

Dictionary

Alice	D1:2	D2:2	D3:2
...			
book	D1:1		
...			
Cheshire Cat	D1:1	D3:1	
...			
grin	D3:2		
...			
King	D1:1	D2:1	
...			
Queen	D2:1	D3:1	
...			
table	D2:1	D3:1	
...			



Inverted Index

Dictionary

Alice	D1:2	D2:2	D3:2
...			

book	D1:1
...	

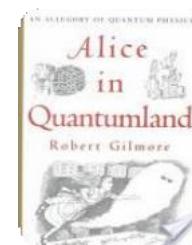
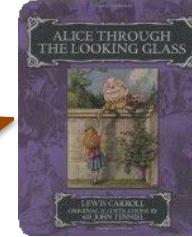
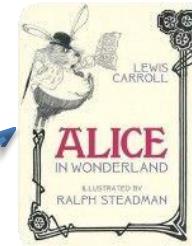
Cheshire Cat	D1:1	D3:1
...		

grin	D3:2
...	

King	D1:1	D2:1
...		

Queen	D2:1	D3:1
...		

table	D2:1	D3:1
...		



Inverted Index

Dictionary

Alice	D1:2	D2:2	D3:2
...			

book	D1:1
...	

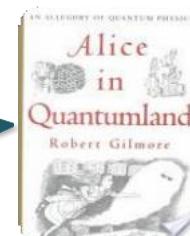
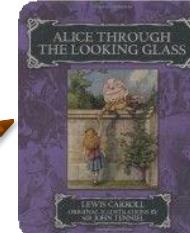
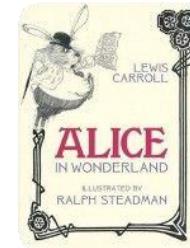
Cheshire Cat	D1:1	D3:1
...		

grin	D3:2
...	

King	D1:1	D2:1
...		

Queen	D2:1	D3:1
...		

table	D2:1	D3:1
...		



Inverted Index

Dictionary

Alice	D1:2	D2:2	D3:2
...			

book	D1:1
...	

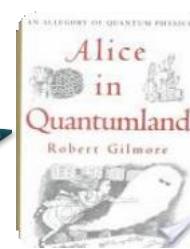
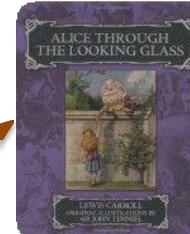
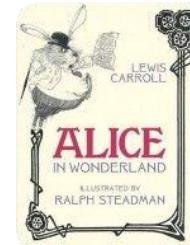
Cheshire Cat	D1:1	D3:1
...		

grin	D3:2
...	

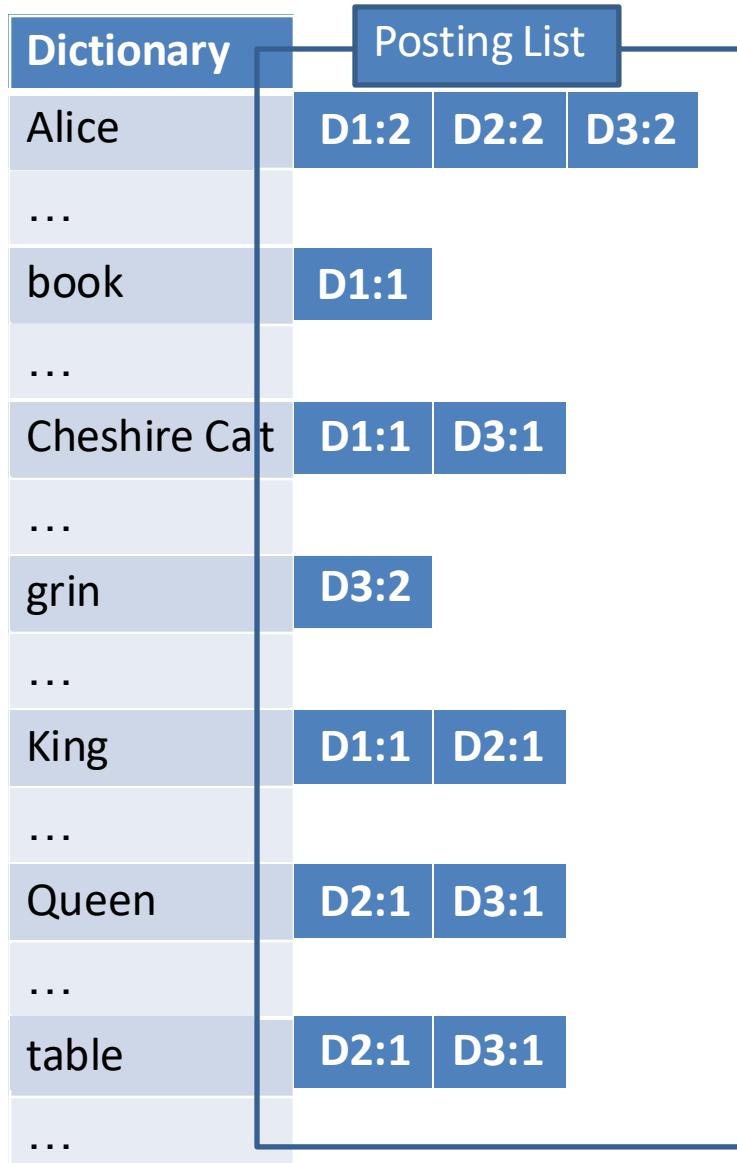
King	D1:1	D2:1
...		

Queen	D2:1	D3:1
...		

table	D2:1	D3:1
...		

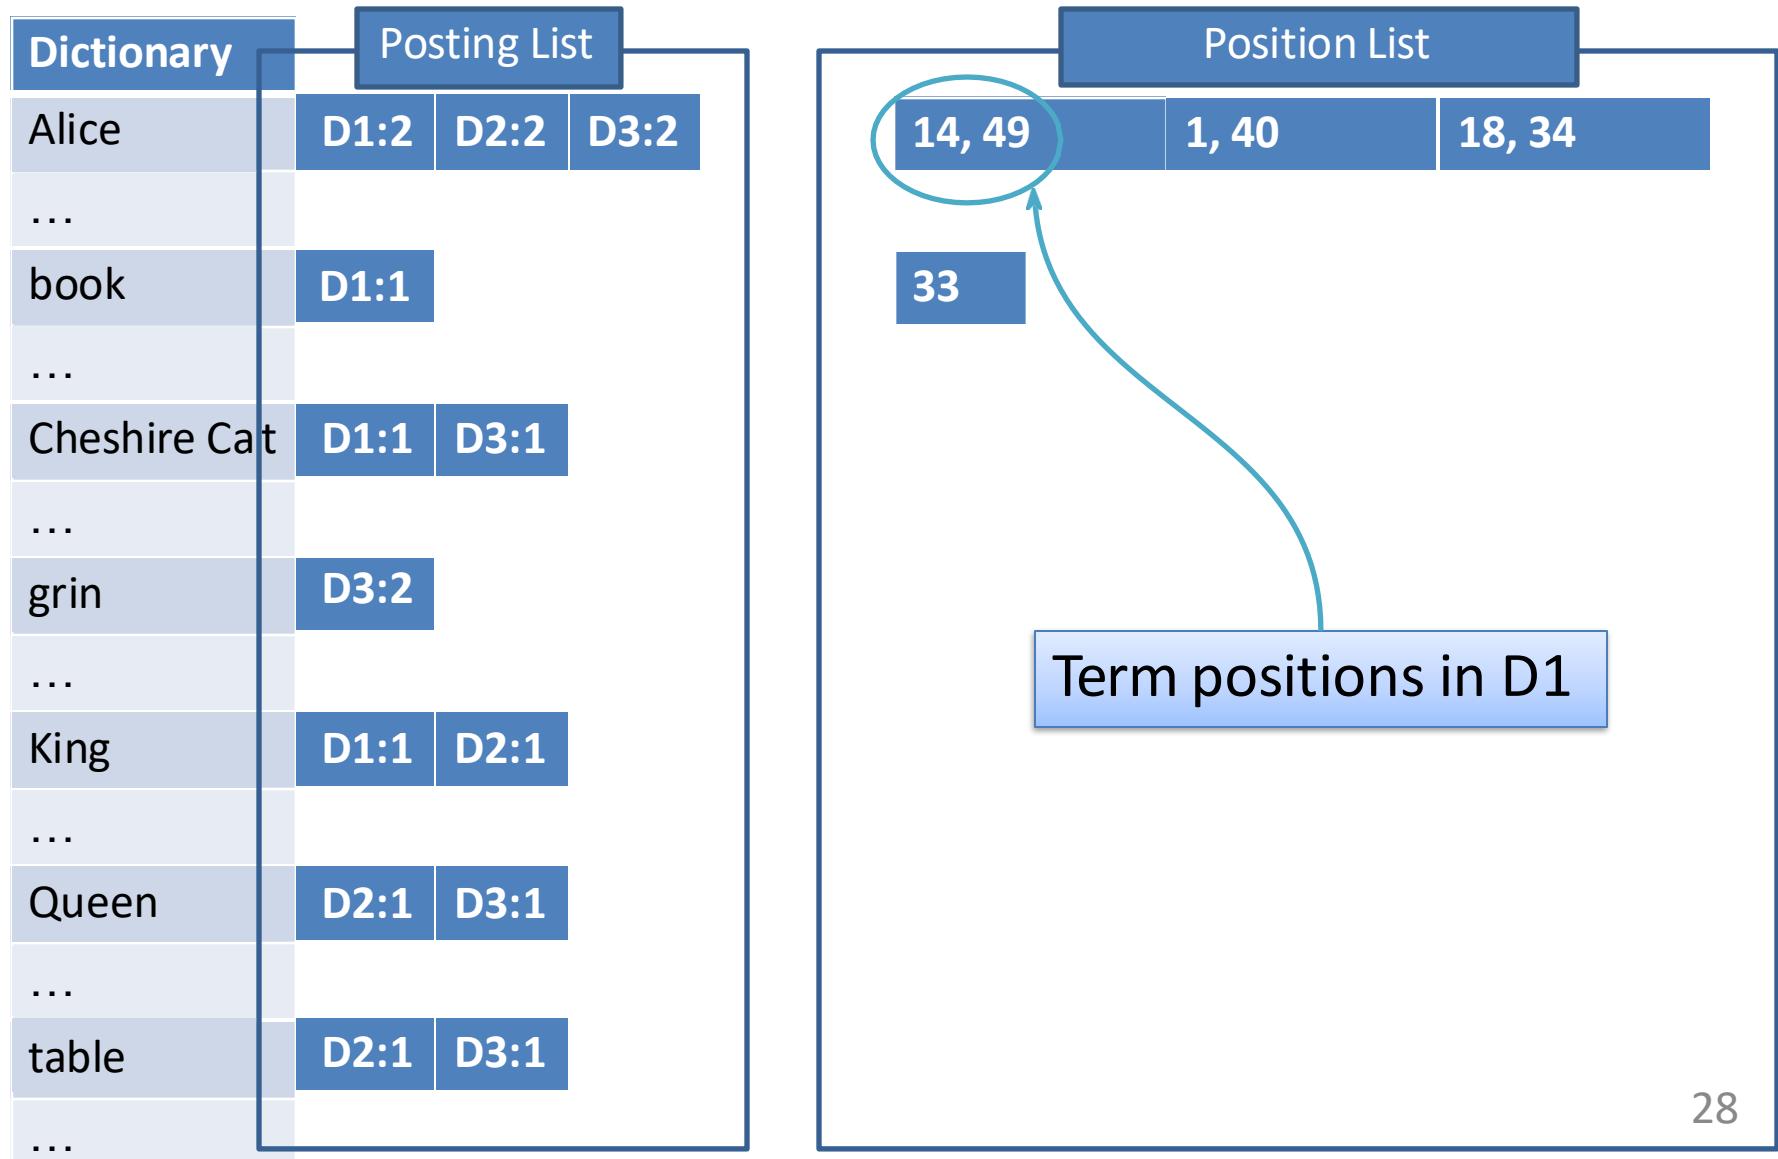


Inverted Index



In more advanced models, we also take into account **the precise position** of the token in the document.

Inverted Index



Inverted Index: query processing

Dictionary	D1:2	D2:2	D3:2
Alice	D1:2	D2:2	D3:2
...			
book	D1:1		
...			
Cheshire Cat	D1:1	D3:1	
...			
grin	D3:2		
...			
King	D1:1	D2:1	
...			
Queen	D2:1	D3:1	
...			
table	D2:1	D3:1	
...			

AND (intersection \cap)

$\langle \text{Alice} \rangle \cap \langle \text{King} \rangle$

Inverted Index: query processing

Dictionary		
Alice	D1:2	D2:2
...		
book	D1:1	
...		
Cheshire Cat	D1:1	D3:1
...		
grin	D3:2	
...		
King	D1:1	D2:1
...		
Queen	D2:1	D3:1
...		
table	D2:1	D3:1
...		

AND (intersection \cap)

D1:2 | D2:2 | D3:2

<Alice> \cap <King>

D1:1 | D2:1

Inverted Index: query processing

Dictionary		
Alice	D1:2	D2:2
...		
book	D1:1	
...		
Cheshire Cat	D1:1	D3:1
...		
grin	D3:2	
...		
King	D1:1	D2:1
...		
Queen	D2:1	D3:1
...		
table	D2:1	D3:1
...		

AND (intersection \cap)

D1:2 | D2:2 | D3:2

<Alice> \cap <King>

D1:1 | D2:1

Alice AND King -> (D1, D2)

Inverted Index: query processing

Dictionary	D1:2	D2:2	D3:2
Alice	D1:2	D2:2	D3:2
...			
book	D1:1		
...			
Cheshire Cat	D1:1	D3:1	
...			
grin	D3:2		
...			
King	D1:1	D2:1	
...			
Queen	D2:1	D3:1	
...			
table	D2:1	D3:1	
...			

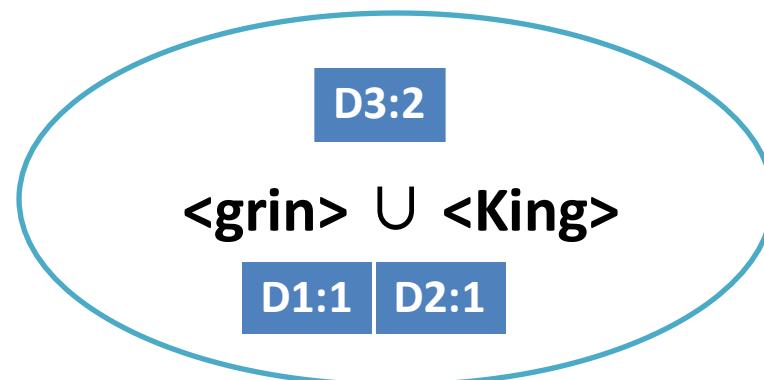
OR (union \cup)

$\langle\text{grin}\rangle \cup \langle\text{King}\rangle$

Inverted Index: query processing

Dictionary			
Alice	D1:2	D2:2	D3:2
...			
book	D1:1		
...			
Cheshire Cat	D1:1	D3:1	
...			
grin	D3:2		
...			
King	D1:1	D2:1	
...			
Queen	D2:1	D3:1	
...			
table	D2:1	D3:1	
...			

OR (union \cup)



$\langle \text{grin} \rangle \cup \langle \text{King} \rangle$

grin OR King -> (D1, D2, D3)

Inverted Index: query processing

Dictionary	D1:2	D2:2	D3:2
Alice	D1:2	D2:2	D3:2
...			
book	D1:1		
...			
Cheshire Cat	D1:1	D3:1	
...			
grin	D3:2		
...			
King	D1:1	D2:1	
...			
Queen	D2:1	D3:1	
...			
table	D2:1	D3:1	
...			

NOT (complement \)

<Alice> \ <grin>

Inverted Index: query processing

Dictionary			
Alice		D1:2	D2:2
		D3:2	
...			
book		D1:1	
...			
Cheshire Cat		D1:1	D3:1
...			
grin		D3:2	
...			
King		D1:1	D2:1
...			
Queen		D2:1	D3:1
...			
table		D2:1	D3:1
...			

NOT (complement \)

D1:2 | D2:2 | D3:2

<Alice> \ <grin>

D3:2

Alice NOT grin -> (D1, D2)

Term-weight

- Measures the term relevance in a document
 - component value in the document representation
 - TF
 - TF (term frequency): term occurrences in the document

Term-weight

- Measures the term relevance in a document
 - component value in the document representation
 - TF
 - TF (term frequency): term occurrences in the document
- Is there any problem with simple counting?

Term-weight

- Measures the term relevance in a document
 - component value in the document representation
 - TF
 - TF (term frequency): term occurrences in the document
- Is there any problem with simple counting?
 - Yes, poorly informative words have high frequency

A toy example

- D1: "Alice is in Wonderland"
- D2: "Alice loves the Cheshire Cat"
- D3: "The Cat is very curious"

A toy example

- D1: "Alice is in Wonderland"
- D2: "Alice loves the Cheshire Cat"
- D3: "The Cat is very curious"

Termini	D1	D2	D3
Alice	1	1	0
is	1	0	1
in	1	0	0
Wonderland	1	0	0
loves	0	1	0
the	0	1	1
Cheshire	0	1	0
Cat	0	1	1
very	0	0	1
curious	0	1	1

A toy example

- D1: "Alice is in Wonderland"
- D2: "Alice loves the Cheshire Cat"
- D3: "The Cat is very curious"

Termini	D1	D2	D3
Alice	1	1	0
is	1	0	1
in	1	0	0
Wonderland	1	0	0
loves	0	1	0
the	0	1	1
Cheshire	0	1	0
Cat	0	1	1
very	0	0	1
curious	0	1	1

It is hard to catch the real 'meaning' of a document.

Possible solutions: (a) remove noise; (b) give more importance to more relevant terms

Term-weight

- Measures the term relevance in a document
 - component value in the document representation
 - TF*IDF
 - TF (term frequency): term occurrences in the document
 - IDF (inverse document frequency): inverse to the number of documents in which the term occurs

$$tf * idf(t, d) = tf(t, d) * \log \frac{|D|}{|\{d \in D : t \in d\}|}$$

Term-weight

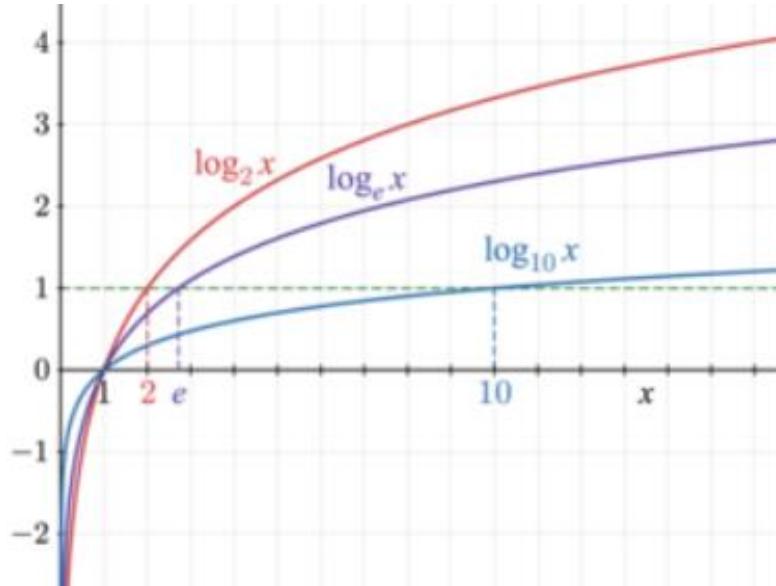
- Measures the term relevance in a document
 - component value in the document representation
 - TF*IDF
 - TF (term frequency): term occurrences in the document
 - IDF (inverse document frequency): inverse to the number of documents in which the term occurs

$$tf * idf(t, d) = tf(t, d) * \log \frac{|D|}{|\{d \in D : t \in d\}|}$$

Number of documents in the collection

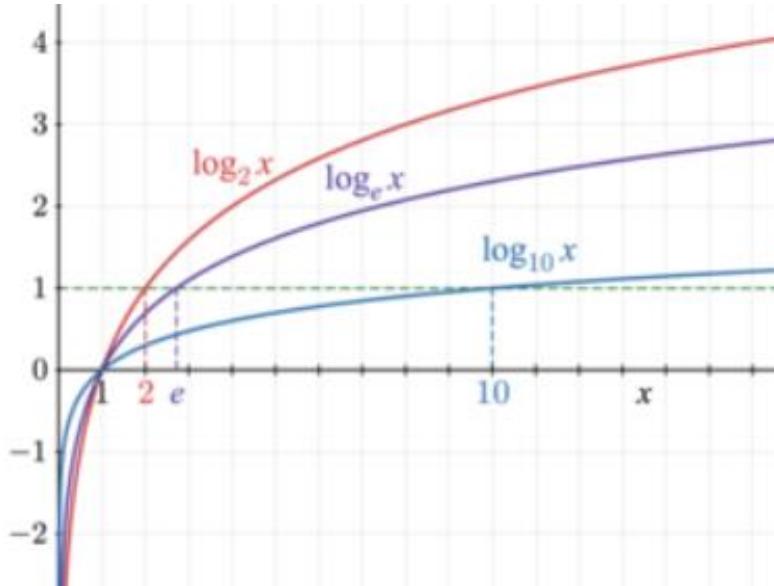
idf

Term-weight



$$tf * idf(t, d) = tf(t, d) * \log \frac{|D|}{|\{d \in D : t \in d\}|}$$

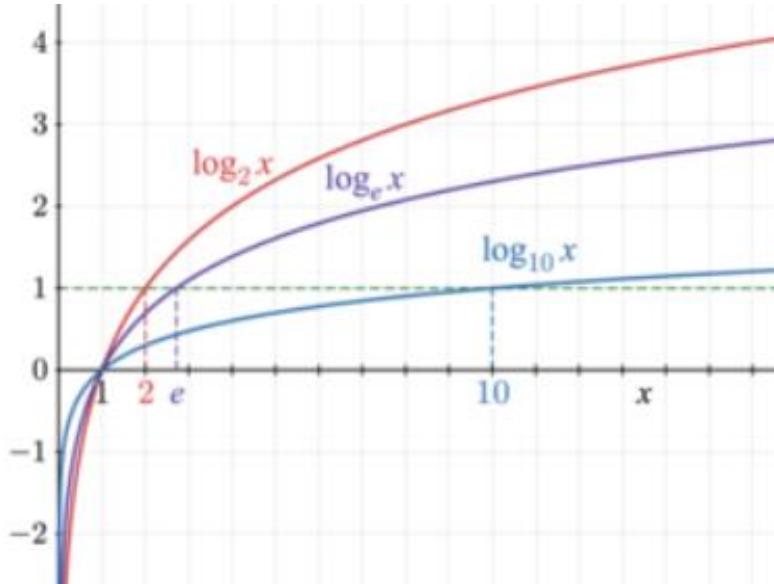
Term-weight



What happens if a term appears in many documents?

$$tf * idf(t, d) = tf(t, d) * \log \frac{|D|}{|\{d \in D : t \in d\}|}$$

Term-weight

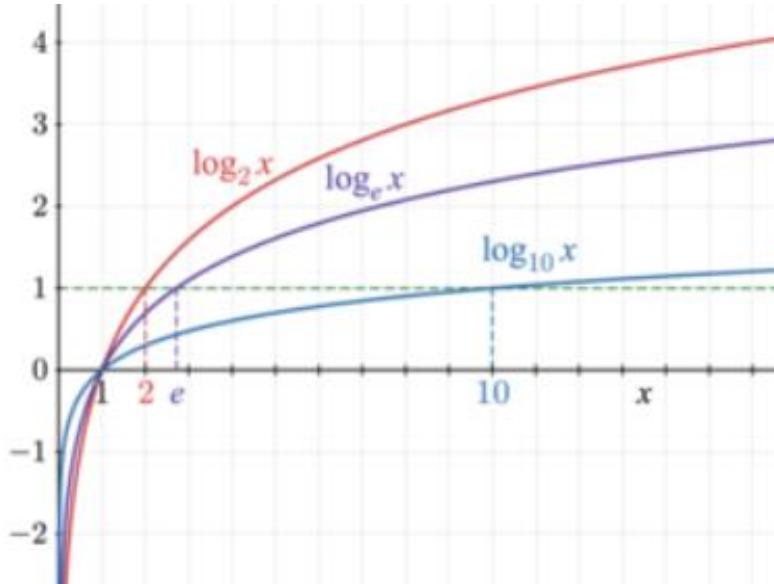


What happens if a term appears in many documents?

The ratio is close to 1, so the logarithm is close to 0

$$tf * idf(t, d) = tf(t, d) * \log \frac{|D|}{|\{d \in D : t \in d\}|}$$

Term-weight



What happens if a term appears in many documents?

The ratio is close to 1, so the logarithm is close to 0

Conversely, if a term appears in just a few document, **the ratio is high and the logarithm is high as well**

$$tf * idf(t, d) = tf(t, d) * \log \frac{|D|}{|\{d \in D : t \in d\}|}$$

TF*IDF insights

- **increases proportionally** to the term frequency in the document
- **decreases** to the number of documents in which the term belongs
 - common words are generally more frequent in the collection
- **IDF depends on the collection, TF on the document**

Inverted index/TF*IDF

- TF: computed by term occurrences in the posting list
- IDF: computed by the posting list cardinality

Alice [D1:2 | D2:2 | D3:2]

$|D|=100$

$$\log \frac{|D|}{|\{d \in D : t \in d\}|}$$

$$tf * idf(Alice, D1) =$$

Inverted index/TF*IDF

- TF: computed by term occurrences in the posting list
- IDF: computed by the posting list cardinality



TF

$$tf * idf(Alice, D1) = 2 * \log \frac{100}{3}$$

Inverted index/TF*IDF

- TF: computed by term occurrences in the posting list
- IDF: computed by the posting list cardinality

The diagram illustrates the calculation of TF-IDF for the term "Alice" in document D1. It shows the inverted index entry for "Alice" (D1:2 | D2:2 | D3:2) and the total number of documents |D|=100. A blue arrow labeled "IDF" points from the term frequency calculation to the IDF component of the formula.

$$tf * idf(Alice, D1) = 2 * \log \frac{100}{3}$$